

Train Object Detector Using R-CNN Deep Learning

César Martín Guijarro & Pablo García García

Firstly, we'll download the Cifar10 dataset and its pretrained convolutional neural network from the given URL.

```
cifar10Data = tempdir;  
url = 'https://www.cs.toronto.edu/~kriz/cifar-10-matlab.tar.gz';  
helperCIFAR10Data.download(url, cifar10Data);
```

Then, we load the data:

```
[trainingImages, trainingLabels, testImages, testLabels] =  
helperCIFAR10Data.load(cifar10Data);  
disp(size(trainingImages));
```

32 32 3 50000

Now we can check the different categories of the images that the dataset contains.

```
numImageCategories = 10;  
disp(categories(trainingLabels));
```

```
{'airplane' }  
{ 'automobile' }  
{ 'bird' }  
{ 'cat' }  
{ 'deer' }  
{ 'dog' }  
{ 'frog' }  
{ 'horse' }  
{ 'ship' }  
{ 'truck' }
```

Let's see some of the images of the dataset:

```
figure;  
thumbnails = trainingImages(:,:,1:100);  
montage(thumbnails);
```



Now we are going to create the CNN. Later, in our slide presentation we will explain how the CNN works.

```
[height, width, numChannels, ~] = size(trainingImages);
imageSize = [height width numChannels];
inputLayer = imageInputLayer(imageSize)
```

```
inputLayer =
  ImageInputLayer with properties:
        Name: ''
        InputSize: [32 32 3]
        SplitComplexInputs: 0
```

Hyperparameters

```
DataAugmentation: 'none'
Normalization: 'zerocenter'
NormalizationDimension: 'auto'
Mean: []
```

```
filterSize = [5 5];
numFilters = 32;

middleLayers = [
    convolution2dLayer(filterSize, numFilters, 'Padding', 2)
    reluLayer()
    maxPooling2dLayer(3, 'Stride', 2)
    convolution2dLayer(filterSize, numFilters, 'Padding', 2)
    reluLayer()
    maxPooling2dLayer(3, 'Stride', 2)
    convolution2dLayer(filterSize, 2 * numFilters, 'Padding', 2)
    reluLayer()
    maxPooling2dLayer(3, 'Stride', 2)
];

finalLayers = [
    fullyConnectedLayer(64)
    reluLayer
    fullyConnectedLayer(numImageCategories)
    softmaxLayer
    classificationLayer
];

layers = [
    inputLayer
    middleLayers
    finalLayers
];

layers(2).Weights = 0.0001 * randn([filterSize numChannels numFilters]);
```

Now we are able to train the net. We'll be using a NVIDIA GPU to make this process faster.

```
opts = trainingOptions('sgdm', ...
    'Momentum', 0.9, ...
    'InitialLearnRate', 0.001, ...
    'LearnRateSchedule', 'piecewise', ...
    'LearnRateDropFactor', 0.1, ...
    'LearnRateDropPeriod', 8, ...
    'L2Regularization', 0.004, ...
    'MaxEpochs', 40, ...
    'MiniBatchSize', 128, ...
    'Verbose', true);

cifar10Net = trainNetwork(trainingImages, trainingLabels, layers, opts);
```

Training on single GPU.

Initializing input data normalization.

| Epoch | Iteration | Time Elapsed (hh:mm:ss) | Mini-batch Accuracy | Mini-batch Loss | Base Learning Rate |
|-------|-----------|----------------------------|------------------------|--------------------|-----------------------|
| 1 | 1 | 00:00:08 | 8.59% | 2.3037 | 0.0010 |
| 1 | 50 | 00:00:09 | 24.22% | 2.0412 | 0.0010 |
| 1 | 100 | 00:00:10 | 28.91% | 1.7514 | 0.0010 |
| 1 | 150 | 00:00:11 | 46.88% | 1.6048 | 0.0010 |
| 1 | 200 | 00:00:12 | 45.31% | 1.6095 | 0.0010 |
| 1 | 250 | 00:00:13 | 34.38% | 1.7327 | 0.0010 |
| 1 | 300 | 00:00:13 | 50.78% | 1.4226 | 0.0010 |
| 1 | 350 | 00:00:14 | 49.22% | 1.3921 | 0.0010 |
| 2 | 400 | 00:00:15 | 35.94% | 1.6275 | 0.0010 |
| 2 | 450 | 00:00:16 | 46.88% | 1.3082 | 0.0010 |
| 2 | 500 | 00:00:17 | 49.22% | 1.4088 | 0.0010 |
| 2 | 550 | 00:00:17 | 59.38% | 1.2249 | 0.0010 |
| 2 | 600 | 00:00:18 | 50.78% | 1.3892 | 0.0010 |
| 2 | 650 | 00:00:19 | 57.03% | 1.2395 | 0.0010 |
| 2 | 700 | 00:00:20 | 49.22% | 1.4112 | 0.0010 |
| 2 | 750 | 00:00:20 | 52.34% | 1.4075 | 0.0010 |
| 3 | 800 | 00:00:21 | 56.25% | 1.2624 | 0.0010 |
| 3 | 850 | 00:00:22 | 57.03% | 1.2232 | 0.0010 |
| 3 | 900 | 00:00:23 | 60.94% | 1.1718 | 0.0010 |
| 3 | 950 | 00:00:23 | 61.72% | 1.0258 | 0.0010 |
| 3 | 1000 | 00:00:24 | 59.38% | 1.1980 | 0.0010 |
| 3 | 1050 | 00:00:25 | 53.12% | 1.2636 | 0.0010 |
| 3 | 1100 | 00:00:26 | 52.34% | 1.3793 | 0.0010 |
| 3 | 1150 | 00:00:26 | 60.16% | 1.0044 | 0.0010 |
| 4 | 1200 | 00:00:27 | 63.28% | 1.0616 | 0.0010 |
| 4 | 1250 | 00:00:28 | 61.72% | 1.2183 | 0.0010 |
| 4 | 1300 | 00:00:29 | 65.62% | 0.9688 | 0.0010 |
| 4 | 1350 | 00:00:29 | 69.53% | 0.9060 | 0.0010 |
| 4 | 1400 | 00:00:30 | 61.72% | 1.0383 | 0.0010 |
| 4 | 1450 | 00:00:31 | 65.62% | 0.9773 | 0.0010 |
| 4 | 1500 | 00:00:32 | 64.84% | 1.1022 | 0.0010 |
| 4 | 1550 | 00:00:32 | 66.41% | 1.0626 | 0.0010 |
| 5 | 1600 | 00:00:33 | 65.62% | 0.9074 | 0.0010 |
| 5 | 1650 | 00:00:34 | 67.19% | 1.0650 | 0.0010 |
| 5 | 1700 | 00:00:35 | 58.59% | 1.1035 | 0.0010 |
| 5 | 1750 | 00:00:35 | 60.94% | 0.9998 | 0.0010 |
| 5 | 1800 | 00:00:36 | 56.25% | 1.1109 | 0.0010 |
| 5 | 1850 | 00:00:37 | 66.41% | 0.9760 | 0.0010 |
| 5 | 1900 | 00:00:38 | 54.69% | 1.2310 | 0.0010 |
| 5 | 1950 | 00:00:38 | 64.06% | 1.0031 | 0.0010 |
| 6 | 2000 | 00:00:39 | 60.94% | 1.1367 | 0.0010 |
| 6 | 2050 | 00:00:40 | 64.84% | 0.8786 | 0.0010 |
| 6 | 2100 | 00:00:40 | 64.84% | 1.0510 | 0.0010 |
| 6 | 2150 | 00:00:41 | 70.31% | 0.9055 | 0.0010 |
| 6 | 2200 | 00:00:42 | 63.28% | 0.9582 | 0.0010 |
| 6 | 2250 | 00:00:43 | 64.06% | 0.9657 | 0.0010 |
| 6 | 2300 | 00:00:43 | 68.75% | 0.8389 | 0.0010 |
| 7 | 2350 | 00:00:44 | 65.62% | 0.9936 | 0.0010 |
| 7 | 2400 | 00:00:45 | 75.00% | 0.7526 | 0.0010 |
| 7 | 2450 | 00:00:46 | 72.66% | 0.8487 | 0.0010 |
| 7 | 2500 | 00:00:46 | 67.97% | 0.8530 | 0.0010 |
| 7 | 2550 | 00:00:47 | 63.28% | 1.0280 | 0.0010 |
| 7 | 2600 | 00:00:48 | 72.66% | 0.8459 | 0.0010 |
| 7 | 2650 | 00:00:48 | 67.19% | 1.0060 | 0.0010 |
| 7 | 2700 | 00:00:49 | 64.84% | 0.9763 | 0.0010 |
| 8 | 2750 | 00:00:50 | 62.50% | 1.0739 | 0.0010 |
| 8 | 2800 | 00:00:51 | 69.53% | 0.9466 | 0.0010 |
| 8 | 2850 | 00:00:51 | 77.34% | 0.7462 | 0.0010 |

| | | | | | |
|----|------|----------|--------|--------|--------|
| 8 | 2900 | 00:00:52 | 73.44% | 0.8096 | 0.0010 |
| 8 | 2950 | 00:00:53 | 68.75% | 0.8545 | 0.0010 |
| 8 | 3000 | 00:00:54 | 75.78% | 0.7059 | 0.0010 |
| 8 | 3050 | 00:00:54 | 66.41% | 1.0119 | 0.0010 |
| 8 | 3100 | 00:00:55 | 75.78% | 0.6426 | 0.0010 |
| 9 | 3150 | 00:00:56 | 75.00% | 0.8256 | 0.0001 |
| 9 | 3200 | 00:00:57 | 71.88% | 0.8012 | 0.0001 |
| 9 | 3250 | 00:00:57 | 78.12% | 0.6230 | 0.0001 |
| 9 | 3300 | 00:00:58 | 75.78% | 0.7534 | 0.0001 |
| 9 | 3350 | 00:00:59 | 73.44% | 0.6201 | 0.0001 |
| 9 | 3400 | 00:01:00 | 78.12% | 0.5866 | 0.0001 |
| 9 | 3450 | 00:01:00 | 78.12% | 0.6060 | 0.0001 |
| 9 | 3500 | 00:01:01 | 82.03% | 0.5479 | 0.0001 |
| 10 | 3550 | 00:01:02 | 77.34% | 0.6012 | 0.0001 |
| 10 | 3600 | 00:01:03 | 73.44% | 0.8093 | 0.0001 |
| 10 | 3650 | 00:01:03 | 77.34% | 0.7267 | 0.0001 |
| 10 | 3700 | 00:01:04 | 72.66% | 0.6763 | 0.0001 |
| 10 | 3750 | 00:01:05 | 75.78% | 0.6993 | 0.0001 |
| 10 | 3800 | 00:01:06 | 76.56% | 0.6848 | 0.0001 |
| 10 | 3850 | 00:01:07 | 75.00% | 0.8749 | 0.0001 |
| 10 | 3900 | 00:01:07 | 78.12% | 0.7080 | 0.0001 |
| 11 | 3950 | 00:01:08 | 72.66% | 0.8838 | 0.0001 |
| 11 | 4000 | 00:01:09 | 75.00% | 0.6583 | 0.0001 |
| 11 | 4050 | 00:01:10 | 74.22% | 0.7222 | 0.0001 |
| 11 | 4100 | 00:01:10 | 77.34% | 0.7884 | 0.0001 |
| 11 | 4150 | 00:01:11 | 70.31% | 0.7545 | 0.0001 |
| 11 | 4200 | 00:01:12 | 71.88% | 0.7922 | 0.0001 |
| 11 | 4250 | 00:01:12 | 78.91% | 0.5906 | 0.0001 |
| 12 | 4300 | 00:01:13 | 68.75% | 0.7395 | 0.0001 |
| 12 | 4350 | 00:01:14 | 76.56% | 0.5622 | 0.0001 |
| 12 | 4400 | 00:01:15 | 78.12% | 0.6263 | 0.0001 |
| 12 | 4450 | 00:01:15 | 78.91% | 0.6004 | 0.0001 |
| 12 | 4500 | 00:01:16 | 67.19% | 0.9125 | 0.0001 |
| 12 | 4550 | 00:01:17 | 78.91% | 0.5927 | 0.0001 |
| 12 | 4600 | 00:01:18 | 75.78% | 0.6726 | 0.0001 |
| 12 | 4650 | 00:01:18 | 74.22% | 0.8036 | 0.0001 |
| 13 | 4700 | 00:01:19 | 74.22% | 0.7746 | 0.0001 |
| 13 | 4750 | 00:01:20 | 74.22% | 0.6979 | 0.0001 |
| 13 | 4800 | 00:01:20 | 79.69% | 0.5740 | 0.0001 |
| 13 | 4850 | 00:01:21 | 78.91% | 0.6090 | 0.0001 |
| 13 | 4900 | 00:01:22 | 71.88% | 0.7117 | 0.0001 |
| 13 | 4950 | 00:01:23 | 82.81% | 0.5817 | 0.0001 |
| 13 | 5000 | 00:01:23 | 71.88% | 0.8572 | 0.0001 |
| 13 | 5050 | 00:01:24 | 85.94% | 0.4858 | 0.0001 |
| 14 | 5100 | 00:01:25 | 78.12% | 0.6789 | 0.0001 |
| 14 | 5150 | 00:01:26 | 72.66% | 0.7179 | 0.0001 |
| 14 | 5200 | 00:01:26 | 81.25% | 0.5365 | 0.0001 |
| 14 | 5250 | 00:01:27 | 78.12% | 0.6634 | 0.0001 |
| 14 | 5300 | 00:01:28 | 78.91% | 0.5539 | 0.0001 |
| 14 | 5350 | 00:01:29 | 78.91% | 0.5589 | 0.0001 |
| 14 | 5400 | 00:01:29 | 79.69% | 0.5655 | 0.0001 |
| 14 | 5450 | 00:01:30 | 84.38% | 0.5278 | 0.0001 |
| 15 | 5500 | 00:01:31 | 79.69% | 0.5470 | 0.0001 |
| 15 | 5550 | 00:01:31 | 75.00% | 0.7370 | 0.0001 |
| 15 | 5600 | 00:01:32 | 78.12% | 0.6769 | 0.0001 |
| 15 | 5650 | 00:01:33 | 75.00% | 0.6318 | 0.0001 |
| 15 | 5700 | 00:01:34 | 78.91% | 0.6598 | 0.0001 |
| 15 | 5750 | 00:01:34 | 77.34% | 0.6373 | 0.0001 |
| 15 | 5800 | 00:01:35 | 78.91% | 0.8177 | 0.0001 |
| 15 | 5850 | 00:01:36 | 78.12% | 0.6712 | 0.0001 |
| 16 | 5900 | 00:01:37 | 71.88% | 0.8234 | 0.0001 |
| 16 | 5950 | 00:01:37 | 79.69% | 0.6079 | 0.0001 |
| 16 | 6000 | 00:01:38 | 75.78% | 0.6762 | 0.0001 |
| 16 | 6050 | 00:01:39 | 78.91% | 0.7349 | 0.0001 |

| | | | | | |
|----|------|----------|--------|--------|------------|
| 16 | 6100 | 00:01:39 | 71.09% | 0.6916 | 0.0001 |
| 16 | 6150 | 00:01:40 | 71.88% | 0.7394 | 0.0001 |
| 16 | 6200 | 00:01:41 | 79.69% | 0.5514 | 0.0001 |
| 17 | 6250 | 00:01:42 | 75.00% | 0.6506 | 1.0000e-05 |
| 17 | 6300 | 00:01:42 | 78.91% | 0.5271 | 1.0000e-05 |
| 17 | 6350 | 00:01:43 | 80.47% | 0.5964 | 1.0000e-05 |
| 17 | 6400 | 00:01:44 | 78.12% | 0.5737 | 1.0000e-05 |
| 17 | 6450 | 00:01:45 | 72.66% | 0.8468 | 1.0000e-05 |
| 17 | 6500 | 00:01:45 | 79.69% | 0.5616 | 1.0000e-05 |
| 17 | 6550 | 00:01:46 | 78.12% | 0.6396 | 1.0000e-05 |
| 17 | 6600 | 00:01:47 | 77.34% | 0.7440 | 1.0000e-05 |
| 18 | 6650 | 00:01:47 | 75.78% | 0.7174 | 1.0000e-05 |
| 18 | 6700 | 00:01:48 | 76.56% | 0.6764 | 1.0000e-05 |
| 18 | 6750 | 00:01:49 | 79.69% | 0.5438 | 1.0000e-05 |
| 18 | 6800 | 00:01:50 | 82.03% | 0.5570 | 1.0000e-05 |
| 18 | 6850 | 00:01:50 | 74.22% | 0.7061 | 1.0000e-05 |
| 18 | 6900 | 00:01:51 | 83.59% | 0.5440 | 1.0000e-05 |
| 18 | 6950 | 00:01:52 | 71.09% | 0.8034 | 1.0000e-05 |
| 18 | 7000 | 00:01:53 | 85.94% | 0.4745 | 1.0000e-05 |
| 19 | 7050 | 00:01:53 | 79.69% | 0.6143 | 1.0000e-05 |
| 19 | 7100 | 00:01:54 | 74.22% | 0.6699 | 1.0000e-05 |
| 19 | 7150 | 00:01:55 | 85.16% | 0.5160 | 1.0000e-05 |
| 19 | 7200 | 00:01:56 | 82.03% | 0.5920 | 1.0000e-05 |
| 19 | 7250 | 00:01:56 | 82.03% | 0.5086 | 1.0000e-05 |
| 19 | 7300 | 00:01:57 | 78.91% | 0.5292 | 1.0000e-05 |
| 19 | 7350 | 00:01:58 | 83.59% | 0.5131 | 1.0000e-05 |
| 19 | 7400 | 00:01:58 | 84.38% | 0.4854 | 1.0000e-05 |
| 20 | 7450 | 00:01:59 | 79.69% | 0.5242 | 1.0000e-05 |
| 20 | 7500 | 00:02:00 | 74.22% | 0.7193 | 1.0000e-05 |
| 20 | 7550 | 00:02:01 | 82.03% | 0.6419 | 1.0000e-05 |
| 20 | 7600 | 00:02:01 | 76.56% | 0.6057 | 1.0000e-05 |
| 20 | 7650 | 00:02:02 | 80.47% | 0.5943 | 1.0000e-05 |
| 20 | 7700 | 00:02:03 | 78.91% | 0.6293 | 1.0000e-05 |
| 20 | 7750 | 00:02:04 | 78.91% | 0.7870 | 1.0000e-05 |
| 20 | 7800 | 00:02:04 | 79.69% | 0.6409 | 1.0000e-05 |
| 21 | 7850 | 00:02:05 | 73.44% | 0.7965 | 1.0000e-05 |
| 21 | 7900 | 00:02:06 | 78.91% | 0.6008 | 1.0000e-05 |
| 21 | 7950 | 00:02:06 | 75.00% | 0.6832 | 1.0000e-05 |
| 21 | 8000 | 00:02:07 | 78.91% | 0.7300 | 1.0000e-05 |
| 21 | 8050 | 00:02:08 | 73.44% | 0.6359 | 1.0000e-05 |
| 21 | 8100 | 00:02:09 | 73.44% | 0.6968 | 1.0000e-05 |
| 21 | 8150 | 00:02:09 | 82.03% | 0.5241 | 1.0000e-05 |
| 22 | 8200 | 00:02:10 | 77.34% | 0.6524 | 1.0000e-05 |
| 22 | 8250 | 00:02:11 | 79.69% | 0.5401 | 1.0000e-05 |
| 22 | 8300 | 00:02:12 | 81.25% | 0.5829 | 1.0000e-05 |
| 22 | 8350 | 00:02:12 | 79.69% | 0.5560 | 1.0000e-05 |
| 22 | 8400 | 00:02:13 | 72.66% | 0.8347 | 1.0000e-05 |
| 22 | 8450 | 00:02:14 | 80.47% | 0.5515 | 1.0000e-05 |
| 22 | 8500 | 00:02:15 | 79.69% | 0.6365 | 1.0000e-05 |
| 22 | 8550 | 00:02:15 | 77.34% | 0.7440 | 1.0000e-05 |
| 23 | 8600 | 00:02:16 | 75.78% | 0.7043 | 1.0000e-05 |
| 23 | 8650 | 00:02:17 | 75.78% | 0.6643 | 1.0000e-05 |
| 23 | 8700 | 00:02:17 | 82.81% | 0.5390 | 1.0000e-05 |
| 23 | 8750 | 00:02:18 | 82.03% | 0.5475 | 1.0000e-05 |
| 23 | 8800 | 00:02:19 | 74.22% | 0.6992 | 1.0000e-05 |
| 23 | 8850 | 00:02:20 | 84.38% | 0.5418 | 1.0000e-05 |
| 23 | 8900 | 00:02:20 | 71.09% | 0.7999 | 1.0000e-05 |
| 23 | 8950 | 00:02:21 | 85.94% | 0.4732 | 1.0000e-05 |
| 24 | 9000 | 00:02:22 | 78.91% | 0.6097 | 1.0000e-05 |
| 24 | 9050 | 00:02:23 | 74.22% | 0.6614 | 1.0000e-05 |
| 24 | 9100 | 00:02:23 | 85.94% | 0.5104 | 1.0000e-05 |
| 24 | 9150 | 00:02:24 | 82.03% | 0.5862 | 1.0000e-05 |
| 24 | 9200 | 00:02:25 | 82.03% | 0.5037 | 1.0000e-05 |
| 24 | 9250 | 00:02:26 | 80.47% | 0.5267 | 1.0000e-05 |


| | | | | | |
|----|-------|----------|--------|--------|------------|
| 24 | 9300 | 00:02:26 | 84.38% | 0.5114 | 1.0000e-05 |
| 24 | 9350 | 00:02:27 | 85.16% | 0.4812 | 1.0000e-05 |
| 25 | 9400 | 00:02:28 | 78.91% | 0.5150 | 1.0000e-06 |
| 25 | 9450 | 00:02:28 | 75.78% | 0.7051 | 1.0000e-06 |
| 25 | 9500 | 00:02:29 | 83.59% | 0.6348 | 1.0000e-06 |
| 25 | 9550 | 00:02:30 | 75.78% | 0.5958 | 1.0000e-06 |
| 25 | 9600 | 00:02:31 | 81.25% | 0.5870 | 1.0000e-06 |
| 25 | 9650 | 00:02:31 | 79.69% | 0.6219 | 1.0000e-06 |
| 25 | 9700 | 00:02:32 | 77.34% | 0.7773 | 1.0000e-06 |
| 25 | 9750 | 00:02:33 | 78.91% | 0.6312 | 1.0000e-06 |
| 26 | 9800 | 00:02:34 | 72.66% | 0.7920 | 1.0000e-06 |
| 26 | 9850 | 00:02:34 | 78.91% | 0.5862 | 1.0000e-06 |
| 26 | 9900 | 00:02:35 | 75.00% | 0.6861 | 1.0000e-06 |
| 26 | 9950 | 00:02:36 | 79.69% | 0.7238 | 1.0000e-06 |
| 26 | 10000 | 00:02:37 | 73.44% | 0.6216 | 1.0000e-06 |
| 26 | 10050 | 00:02:37 | 73.44% | 0.6845 | 1.0000e-06 |
| 26 | 10100 | 00:02:38 | 82.03% | 0.5174 | 1.0000e-06 |
| 27 | 10150 | 00:02:39 | 78.12% | 0.6520 | 1.0000e-06 |
| 27 | 10200 | 00:02:40 | 79.69% | 0.5348 | 1.0000e-06 |
| 27 | 10250 | 00:02:40 | 80.47% | 0.5798 | 1.0000e-06 |
| 27 | 10300 | 00:02:41 | 80.47% | 0.5488 | 1.0000e-06 |
| 27 | 10350 | 00:02:42 | 72.66% | 0.8264 | 1.0000e-06 |
| 27 | 10400 | 00:02:42 | 81.25% | 0.5435 | 1.0000e-06 |
| 27 | 10450 | 00:02:43 | 78.91% | 0.6291 | 1.0000e-06 |
| 27 | 10500 | 00:02:44 | 77.34% | 0.7377 | 1.0000e-06 |
| 28 | 10550 | 00:02:45 | 75.00% | 0.6986 | 1.0000e-06 |
| 28 | 10600 | 00:02:45 | 77.34% | 0.6585 | 1.0000e-06 |
| 28 | 10650 | 00:02:46 | 84.38% | 0.5394 | 1.0000e-06 |
| 28 | 10700 | 00:02:47 | 82.03% | 0.5417 | 1.0000e-06 |
| 28 | 10750 | 00:02:48 | 73.44% | 0.6997 | 1.0000e-06 |
| 28 | 10800 | 00:02:48 | 84.38% | 0.5357 | 1.0000e-06 |
| 28 | 10850 | 00:02:49 | 73.44% | 0.7938 | 1.0000e-06 |
| 28 | 10900 | 00:02:50 | 85.94% | 0.4696 | 1.0000e-06 |
| 29 | 10950 | 00:02:51 | 78.91% | 0.6088 | 1.0000e-06 |
| 29 | 11000 | 00:02:51 | 74.22% | 0.6581 | 1.0000e-06 |
| 29 | 11050 | 00:02:52 | 85.16% | 0.5073 | 1.0000e-06 |
| 29 | 11100 | 00:02:53 | 82.03% | 0.5782 | 1.0000e-06 |
| 29 | 11150 | 00:02:53 | 81.25% | 0.5024 | 1.0000e-06 |
| 29 | 11200 | 00:02:54 | 77.34% | 0.5271 | 1.0000e-06 |
| 29 | 11250 | 00:02:55 | 84.38% | 0.5057 | 1.0000e-06 |
| 29 | 11300 | 00:02:56 | 87.50% | 0.4766 | 1.0000e-06 |
| 30 | 11350 | 00:02:56 | 79.69% | 0.5152 | 1.0000e-06 |
| 30 | 11400 | 00:02:57 | 75.00% | 0.7047 | 1.0000e-06 |
| 30 | 11450 | 00:02:58 | 82.81% | 0.6316 | 1.0000e-06 |
| 30 | 11500 | 00:02:59 | 75.78% | 0.5971 | 1.0000e-06 |
| 30 | 11550 | 00:02:59 | 81.25% | 0.5840 | 1.0000e-06 |
| 30 | 11600 | 00:03:00 | 78.91% | 0.6208 | 1.0000e-06 |
| 30 | 11650 | 00:03:01 | 78.12% | 0.7774 | 1.0000e-06 |
| 30 | 11700 | 00:03:01 | 78.91% | 0.6318 | 1.0000e-06 |
| 31 | 11750 | 00:03:02 | 72.66% | 0.7893 | 1.0000e-06 |
| 31 | 11800 | 00:03:03 | 78.91% | 0.5859 | 1.0000e-06 |
| 31 | 11850 | 00:03:04 | 75.00% | 0.6850 | 1.0000e-06 |
| 31 | 11900 | 00:03:04 | 79.69% | 0.7213 | 1.0000e-06 |
| 31 | 11950 | 00:03:05 | 73.44% | 0.6209 | 1.0000e-06 |
| 31 | 12000 | 00:03:06 | 73.44% | 0.6839 | 1.0000e-06 |
| 31 | 12050 | 00:03:07 | 82.03% | 0.5178 | 1.0000e-06 |
| 32 | 12100 | 00:03:07 | 78.12% | 0.6516 | 1.0000e-06 |
| 32 | 12150 | 00:03:08 | 80.47% | 0.5348 | 1.0000e-06 |
| 32 | 12200 | 00:03:09 | 80.47% | 0.5784 | 1.0000e-06 |
| 32 | 12250 | 00:03:10 | 80.47% | 0.5488 | 1.0000e-06 |
| 32 | 12300 | 00:03:10 | 72.66% | 0.8255 | 1.0000e-06 |
| 32 | 12350 | 00:03:11 | 81.25% | 0.5429 | 1.0000e-06 |
| 32 | 12400 | 00:03:12 | 78.91% | 0.6289 | 1.0000e-06 |
| 32 | 12450 | 00:03:12 | 77.34% | 0.7376 | 1.0000e-06 |

| | | | | | |
|----|-------|----------|--------|--------|------------|
| 33 | 12500 | 00:03:13 | 75.78% | 0.6972 | 1.0000e-07 |
| 33 | 12550 | 00:03:14 | 76.56% | 0.6580 | 1.0000e-07 |
| 33 | 12600 | 00:03:15 | 84.38% | 0.5390 | 1.0000e-07 |
| 33 | 12650 | 00:03:15 | 82.03% | 0.5415 | 1.0000e-07 |
| 33 | 12700 | 00:03:16 | 73.44% | 0.6985 | 1.0000e-07 |
| 33 | 12750 | 00:03:17 | 85.16% | 0.5355 | 1.0000e-07 |
| 33 | 12800 | 00:03:18 | 73.44% | 0.7932 | 1.0000e-07 |
| 33 | 12850 | 00:03:18 | 86.72% | 0.4687 | 1.0000e-07 |
| 34 | 12900 | 00:03:19 | 78.91% | 0.6083 | 1.0000e-07 |
| 34 | 12950 | 00:03:20 | 74.22% | 0.6582 | 1.0000e-07 |
| 34 | 13000 | 00:03:20 | 85.16% | 0.5070 | 1.0000e-07 |
| 34 | 13050 | 00:03:21 | 82.03% | 0.5777 | 1.0000e-07 |
| 34 | 13100 | 00:03:22 | 81.25% | 0.5017 | 1.0000e-07 |
| 34 | 13150 | 00:03:23 | 77.34% | 0.5265 | 1.0000e-07 |
| 34 | 13200 | 00:03:23 | 84.38% | 0.5058 | 1.0000e-07 |
| 34 | 13250 | 00:03:24 | 87.50% | 0.4764 | 1.0000e-07 |
| 35 | 13300 | 00:03:25 | 79.69% | 0.5143 | 1.0000e-07 |
| 35 | 13350 | 00:03:26 | 75.00% | 0.7027 | 1.0000e-07 |
| 35 | 13400 | 00:03:26 | 82.81% | 0.6307 | 1.0000e-07 |
| 35 | 13450 | 00:03:27 | 75.78% | 0.5967 | 1.0000e-07 |
| 35 | 13500 | 00:03:28 | 82.03% | 0.5825 | 1.0000e-07 |
| 35 | 13550 | 00:03:29 | 78.91% | 0.6188 | 1.0000e-07 |
| 35 | 13600 | 00:03:29 | 78.12% | 0.7762 | 1.0000e-07 |
| 35 | 13650 | 00:03:30 | 78.91% | 0.6312 | 1.0000e-07 |
| 36 | 13700 | 00:03:31 | 72.66% | 0.7890 | 1.0000e-07 |
| 36 | 13750 | 00:03:32 | 79.69% | 0.5846 | 1.0000e-07 |
| 36 | 13800 | 00:03:32 | 75.00% | 0.6850 | 1.0000e-07 |
| 36 | 13850 | 00:03:33 | 79.69% | 0.7214 | 1.0000e-07 |
| 36 | 13900 | 00:03:34 | 73.44% | 0.6193 | 1.0000e-07 |
| 36 | 13950 | 00:03:34 | 73.44% | 0.6829 | 1.0000e-07 |
| 36 | 14000 | 00:03:35 | 82.03% | 0.5178 | 1.0000e-07 |
| 37 | 14050 | 00:03:36 | 78.12% | 0.6512 | 1.0000e-07 |
| 37 | 14100 | 00:03:37 | 80.47% | 0.5347 | 1.0000e-07 |
| 37 | 14150 | 00:03:37 | 80.47% | 0.5785 | 1.0000e-07 |
| 37 | 14200 | 00:03:38 | 80.47% | 0.5474 | 1.0000e-07 |
| 37 | 14250 | 00:03:39 | 72.66% | 0.8246 | 1.0000e-07 |
| 37 | 14300 | 00:03:40 | 81.25% | 0.5415 | 1.0000e-07 |
| 37 | 14350 | 00:03:40 | 78.91% | 0.6285 | 1.0000e-07 |
| 37 | 14400 | 00:03:41 | 77.34% | 0.7370 | 1.0000e-07 |
| 38 | 14450 | 00:03:42 | 75.78% | 0.6969 | 1.0000e-07 |
| 38 | 14500 | 00:03:42 | 76.56% | 0.6577 | 1.0000e-07 |
| 38 | 14550 | 00:03:43 | 84.38% | 0.5390 | 1.0000e-07 |
| 38 | 14600 | 00:03:44 | 82.03% | 0.5413 | 1.0000e-07 |
| 38 | 14650 | 00:03:45 | 73.44% | 0.6984 | 1.0000e-07 |
| 38 | 14700 | 00:03:45 | 85.16% | 0.5353 | 1.0000e-07 |
| 38 | 14750 | 00:03:46 | 73.44% | 0.7933 | 1.0000e-07 |
| 38 | 14800 | 00:03:47 | 86.72% | 0.4686 | 1.0000e-07 |
| 39 | 14850 | 00:03:48 | 78.91% | 0.6080 | 1.0000e-07 |
| 39 | 14900 | 00:03:48 | 74.22% | 0.6581 | 1.0000e-07 |
| 39 | 14950 | 00:03:49 | 85.16% | 0.5070 | 1.0000e-07 |
| 39 | 15000 | 00:03:50 | 82.03% | 0.5777 | 1.0000e-07 |
| 39 | 15050 | 00:03:50 | 81.25% | 0.5018 | 1.0000e-07 |
| 39 | 15100 | 00:03:51 | 77.34% | 0.5267 | 1.0000e-07 |
| 39 | 15150 | 00:03:52 | 84.38% | 0.5058 | 1.0000e-07 |
| 39 | 15200 | 00:03:53 | 87.50% | 0.4765 | 1.0000e-07 |
| 40 | 15250 | 00:03:53 | 79.69% | 0.5143 | 1.0000e-07 |
| 40 | 15300 | 00:03:54 | 75.00% | 0.7027 | 1.0000e-07 |
| 40 | 15350 | 00:03:55 | 82.81% | 0.6304 | 1.0000e-07 |
| 40 | 15400 | 00:03:56 | 75.78% | 0.5968 | 1.0000e-07 |
| 40 | 15450 | 00:03:56 | 82.03% | 0.5824 | 1.0000e-07 |
| 40 | 15500 | 00:03:57 | 79.69% | 0.6189 | 1.0000e-07 |
| 40 | 15550 | 00:03:58 | 78.91% | 0.7765 | 1.0000e-07 |
| 40 | 15600 | 00:03:58 | 78.91% | 0.6312 | 1.0000e-07 |


```
|=====|  
Training finished: Max epochs completed.
```

Once the training has finished, we can see what is happening in some of the middle layers. In this case, it seems as if its trying to detect borders.

```
w = cifar10Net.Layers(2).Weights;  
w = rescale(w);  
  
figure  
montage(w);
```



Now we will check if the network has been trained correctly:

```
YTest = classify(cifar10Net, testImages);
accuracy = sum(YTest == testLabels)/numel(testLabels)
```

```
accuracy = 0.7072
```

Detecting STOP signs

Having correctly trained our CNN, we will try to use it to detect STOP signs. To do so, we need to load only the images we will use. In this case will follow the MATLAB example.

```
data = load('stopSignsAndCars.mat', 'stopSignsAndCars');
stopSignsAndCars = data.stopSignsAndCars;

visiondata = fullfile(toolboxdir('vision'), 'visiondata');
stopSignsAndCars.imageFilename = fullfile(visiondata,
stopSignsAndCars.imageFilename);

summary(stopSignsAndCars)
```

Variables:

```
imageFilename: 41x1 cell array of character vectors
```

```
stopSign: 41x1 cell
```

```
carRear: 41x1 cell
```

```
carFront: 41x1 cell
```

In this dataset, every object of interest has a bounding box surrounding it, so to detect stop signs we need to load only the images which them. Let's see an example of an image with an stop sign with its boundary box.

```
stopSigns = stopSignsAndCars(:, {'imageFilename','stopSign'});

I = imread(stopSigns.imageFilename{1});
I = insertObjectAnnotation(I,'Rectangle',stopSigns.stopSign{1},'stop
sign','LineWidth',8);

figure
imshow(I)
```



Now we will train our CNN but not from 0 knowledge, it will be a training that modifies the knowledge of the CIFAR-10 network.

```
options = trainingOptions('sgdm', ...
    'MiniBatchSize', 128, ...
    'InitialLearnRate', 1e-3, ...
    'LearnRateSchedule', 'piecewise', ...
    'LearnRateDropFactor', 0.1, ...
    'LearnRateDropPeriod', 100, ...
    'MaxEpochs', 100, ...
    'Verbose', true);

rcnn = trainRCNNObjectDetector(stopSigns, cifar10Net, options,
    'NegativeOverlapRange', [0 0.3], 'PositiveOverlapRange',[0.5 1]);
```

Starting parallel pool (parpool) using the 'Processes' profile ...
Parallel pool using the 'Processes' profile is shutting down.

Training an R-CNN Object Detector for the following object classes:

* stopSign

--> Extracting region proposals from 41 training images...done.

--> Training a neural network to classify objects in training data...

Training on single GPU.
Initializing input data normalization.

| Epoch | Iteration | Time Elapsed (hh:mm:ss) | Mini-batch Accuracy | Mini-batch Loss | Base Learning Rate |
|-------|-----------|----------------------------|------------------------|--------------------|-----------------------|
| 1 | 1 | 00:00:00 | 57.03% | 0.7087 | 0.0010 |
| 6 | 50 | 00:00:19 | 99.22% | 0.0135 | 0.0010 |
| 12 | 100 | 00:00:38 | 100.00% | 0.0125 | 0.0010 |
| 17 | 150 | 00:00:56 | 100.00% | 0.0010 | 0.0010 |
| 23 | 200 | 00:01:14 | 99.22% | 0.0068 | 0.0010 |
| 28 | 250 | 00:01:32 | 100.00% | 0.0005 | 0.0010 |
| 34 | 300 | 00:01:51 | 100.00% | 0.0009 | 0.0010 |
| 39 | 350 | 00:02:09 | 100.00% | 9.9973e-05 | 0.0010 |
| 45 | 400 | 00:02:27 | 100.00% | 8.3369e-05 | 0.0010 |
| 50 | 450 | 00:02:46 | 100.00% | 0.0004 | 0.0010 |
| 56 | 500 | 00:03:04 | 100.00% | 0.0005 | 0.0010 |
| 62 | 550 | 00:03:23 | 100.00% | 0.0009 | 0.0010 |
| 67 | 600 | 00:03:41 | 100.00% | 0.0003 | 0.0010 |
| 73 | 650 | 00:03:59 | 100.00% | 0.0001 | 0.0010 |
| 78 | 700 | 00:04:18 | 100.00% | 0.0001 | 0.0010 |
| 84 | 750 | 00:04:36 | 100.00% | 0.0003 | 0.0010 |
| 89 | 800 | 00:04:54 | 100.00% | 0.0001 | 0.0010 |
| 95 | 850 | 00:05:13 | 100.00% | 0.0002 | 0.0010 |
| 100 | 900 | 00:05:31 | 100.00% | 0.0001 | 0.0010 |

Training finished: Max epochs completed.

Network training complete.

--> Training bounding box regression models for each object class...100.00%...done.

Detector training complete.

As we've done before, once the training has finished, we can try the net.

```
testImage = imread('stopSignTest.jpg');
[bboxes,score,label] = detect(rcnn, testImage, 'MiniBatchSize', 128);

[score, idx] = max(score);

bbox = bboxes(idx, :);
annotation = sprintf('%s: (Confidence = %f)', label(idx), score);

outputImage = insertObjectAnnotation(testImage, 'rectangle', bbox, annotation);

figure;
imshow(outputImage);
```



Now we can try the model with some photos taken by us:

```
testImage = imread('imagenes_jpg/stop/1.jpg');  
[bboxes, score, label] = detect(rcnn, testImage, 'MiniBatchSize', 128);  
[score, idx] = max(score);  
bbox = bboxes(idx, :);  
annotation = sprintf('%s: (Confidence = %f)', label(idx), score);  
outputImage = insertObjectAnnotation(testImage, 'rectangle', bbox, annotation);  
figure;  
imshow(outputImage);
```




DETECTION OF SIGNS IN OUR OWN PHOTOS

To cover this section we went out to the street and took pictures of different traffic signs, so we can train different CNNs over the Cifar one. We have also made a MATLAB script which generates a dataset with each type of image. In this first example we will try to detect crosswalk signs. To begin with, we show one of our dataset images.

```
load('./imagenes_jpg/peatones/dataset_peatones.mat');  
dataset_pedestrian = dataset;  
  
I_pedestrian = imread(dataset_pedestrian.paths{1});  
I_pedestrian = insertObjectAnnotation(I_pedestrian, 'Rectangle',  
dataset_pedestrian.peatones{1}, 'Pedestrian sign', 'LineWidth', 8);  
  
figure;  
imshow(I_pedestrian);
```



Once we've seen an example image, we can train the net. In this case, we will implement a different algorithm than the propussed one, we will use Adam which will be also explained in our slides. Also we will use our NVIDIA GPU to train faster.

```
options = trainingOptions('adam', ...
    'MiniBatchSize', 128, ...
    'InitialLearnRate', 1e-3, ...
    'LearnRateSchedule', 'piecewise', ...
    'LearnRateDropFactor', 0.1, ...
    'LearnRateDropPeriod', 100, ...
    'MaxEpochs', 100, ...
    'Verbose', true, ...
    'ExecutionEnvironment', 'gpu');

rcnn_pedestrian = trainRCNNObjectDetector(dataset_pedestrian, cifar10Net, options,
    'NegativeOverlapRange', [0 0.3], 'PositiveOverlapRange', [0.5 1]);
```

Starting parallel pool (parpool) using the 'Processes' profile ...
 Parallel pool using the 'Processes' profile is shutting down.

Training an R-CNN Object Detector for the following object classes:

```
* rois
```

```
--> Extracting region proposals from 33 training images...done.
```

```
--> Training a neural network to classify objects in training data...
```

Initializing input data normalization.

| Epoch | Iteration | Time Elapsed (hh:mm:ss) | Mini-batch Accuracy | Mini-batch Loss | Base Learning Rate |
|-------|-----------|----------------------------|------------------------|--------------------|-----------------------|
| 1 | 1 | 00:00:00 | 31.25% | 0.7435 | 0.0010 |
| 6 | 50 | 00:00:12 | 94.53% | 0.1587 | 0.0010 |
| 12 | 100 | 00:00:25 | 100.00% | 0.0002 | 0.0010 |
| 17 | 150 | 00:00:37 | 99.22% | 0.0070 | 0.0010 |
| 23 | 200 | 00:00:49 | 100.00% | 6.2246e-05 | 0.0010 |
| 28 | 250 | 00:01:01 | 100.00% | 0.0074 | 0.0010 |
| 34 | 300 | 00:01:14 | 100.00% | 0.0014 | 0.0010 |
| 39 | 350 | 00:01:26 | 99.22% | 0.0275 | 0.0010 |
| 45 | 400 | 00:01:38 | 100.00% | 0.0003 | 0.0010 |
| 50 | 450 | 00:01:50 | 100.00% | 6.1628e-05 | 0.0010 |
| 56 | 500 | 00:02:02 | 99.22% | 0.0100 | 0.0010 |
| 62 | 550 | 00:02:14 | 100.00% | 0.0017 | 0.0010 |
| 67 | 600 | 00:02:26 | 99.22% | 0.0890 | 0.0010 |
| 73 | 650 | 00:02:38 | 100.00% | 2.2494e-05 | 0.0010 |
| 78 | 700 | 00:02:51 | 100.00% | 0.0023 | 0.0010 |
| 84 | 750 | 00:03:03 | 100.00% | 0.0001 | 0.0010 |
| 89 | 800 | 00:03:15 | 100.00% | 0.0029 | 0.0010 |
| 95 | 850 | 00:03:27 | 100.00% | 8.0533e-05 | 0.0010 |
| 100 | 900 | 00:03:39 | 100.00% | 6.2596e-06 | 0.0010 |

Training finished: Max epochs completed.

Network training complete.

```
--> Training bounding box regression models for each object class...100.00%...done.
```

Detector training complete.

When the model has finished training, we can try with another photo

```
test_image_pedestrian = imread('./imagenes_maps/1.png');
[bboxes, score, label] = detect(rcnn_pedestrian, test_image_pedestrian,
'MiniBatchSize', 128);

[score, idx] = max(score);
bbox = bboxes(idx, :);
annotation = sprintf('%s: (Confidence = %f)', 'Pedestrian sign', score);
outputImage = insertObjectAnnotation(test_image_pedestrian, 'rectangle', bbox,
annotation);
figure;
imshow(outputImage);
```




Also, we modified the recognition code so if more than one object appears in the image, it is also detected.

```
test_image_pedestrian = imread('./imagenes_maps/11.png');
[bboxes, score, label] = detect(rcnn_pedestrian, test_image_pedestrian,
'MiniBatchSize', 128);

instances = 5; %specify how many objects to detect

objects = min(instances, size(bboxes, 1));
[scores_sorted, idx_sorted] = sort(score, 'descend');
outputImage = test_image_pedestrian;

for i = 1:objects
    bbox = bboxes(idx_sorted(i), :);
    annotation = sprintf('%s: (Confidence = %f)', 'Pedestrian sign',
scores_sorted(i));
```

```

    outputImage = insertObjectAnnotation(outputImage, 'rectangle', bbox,
annotation);
end

figure;
imshow(outputImage);

```



Maybe a too many objects are detected, when they shouldn't. To fix that we can set a confidence threshold and show only those whose confidence threshold is bigger than the set one. Here is the alternative version with a minimum score threshold:

```

test_image_pedestrian = imread('./imagenes_jpg/direccion_prohibida/3.jpg');
[bboxes, score, label] = detect(rcnn_pedestrian, test_image_pedestrian,
'MiniBatchSize', 128);

% Set the minimum score threshold
min_score_threshold = 0.7;

% Filter out detections below the threshold
valid_indices = find(score >= min_score_threshold);
bboxes = bboxes(valid_indices, :);
score = score(valid_indices);
label = label(valid_indices);

objects = numel(valid_indices);

```



```

% Sort the scores
[scores_sorted, idx_sorted] = sort(score, 'descend');
outputImage = test_image_pedestrian;

for i = 1:objects
    bbox = bboxes(idx_sorted(i), :);
    annotation = sprintf('%s: (Confidence = %f)', 'Pedestrian sign',
scores_sorted(i));
    outputImage = insertObjectAnnotation(outputImage, 'rectangle', bbox,
annotation);
end

figure;
imshow(outputImage);

```



Heatmap

As shown by the Matlab example, this can also be used to process the entire image and gain information about what the Network is actually seeing.

```
% The trained network is stored within the R-CNN detector
rcnn_pedestrian.Network
```

```
ans =
  SeriesNetwork with properties:

    Layers: [15x1 nnet.cnn.layer.Layer]
  InputNames: {'imageinput'}
  OutputNames: {'rcnnClassification'}
```

As a way to debug, we can extract the activations of the networks. The information we need is stored in different dimensions.

```
test_image_pedestrian = imread('./imagenes_maps/9.png');
featureMap = activations(rcnn_pedestrian.Network, test_image_pedestrian, 14);

% The softmax activations are stored in a 3-D array.
size(featureMap)
```

```
ans = 1x3
    107    149     2
```

In this case the pedestrian signs are stored in the third dimension.

```
rcnn_pedestrian.ClassNames
```

The pedestrian sign feature map is stored in the first channel.


```
pedestrianSignMap = featureMap(:, :, 1);
```

Even though this is a "rough approximation" of what the Network is seeing, it can

```
% Resize stopSignMap for visualization
[height, width, ~] = size(test_image_pedestrian);
pedestrianSignMap = imresize(pedestrianSignMap, [height, width]);

% Visualize the feature map superimposed on the test image.
featureMapOnImage = imfuse(test_image_pedestrian, pedestrianSignMap);

figure
imshow(featureMapOnImage)
```



BAD EXAMPLE: TRAFFIC LIGHT DETECTION

Not everything is perfect, CNNs not always work perfectly. Detecting 3D objects with strange backgrounds is a very hard task. To show that, we have also taken as many photos of traffic lights as of pedestrian signs (more or less), and despite the fact that for pedestrian signs it worked perfectly, now we will see that it is not the case with traffic lights. Maybe it is caused by the fact that photos can be taken in different angles, now what we want to detect is a 3D object and its appearance depends on the angle the photo is taken, not as the sign, independently where you take the photo from, it will always look like a blue square with a white triangle inside. Also the backgrounds of the traffic lights didn't help.

```
load('./imagenes_jpg/semaforo/dataset_semaforo.mat');  
dataset_traffic_lights = dataset;  
  
I_traffic_lights = imread(dataset_traffic_lights.paths{1});  
I_traffic_lights = insertObjectAnnotation(I_traffic_lights, 'Rectangle',  
dataset_traffic_lights.semaforo{1}, 'Traffic light', 'LineWidth', 8);  
  
figure;  
imshow(I_traffic_lights);
```




Once we've seen an example image, we can train the net, also with Adam and our NVIDIA GPU to finish the training earlier.

```
rcnn_traffic_lights = trainRCNNObjectDetector(dataset_traffic_lights, cifar10Net,
options, 'NegativeOverlapRange', [0 0.3], 'PositiveOverlapRange',[0.5 1]);
```

Starting parallel pool (parpool) using the 'Processes' profile ...

Connected to parallel pool with 6 workers.

Training an R-CNN Object Detector for the following object classes:

* rois

--> Extracting region proposals from 31 training images...done.

--> Training a neural network to classify objects in training data...

Initializing input data normalization.

| Epoch | Iteration | Time Elapsed (hh:mm:ss) | Mini-batch Accuracy | Mini-batch Loss | Base Learning Rate |
|-------|-----------|----------------------------|------------------------|--------------------|-----------------------|
| 1 | 1 | 00:00:01 | 64.84% | 0.6583 | 0.0010 |
| 25 | 50 | 00:00:06 | 100.00% | 0.0109 | 0.0010 |

| | | | | | |
|-----|-----|----------|---------|--------|--------|
| 50 | 100 | 00:00:10 | 100.00% | 0.0008 | 0.0010 |
| 75 | 150 | 00:00:15 | 100.00% | 0.0007 | 0.0010 |
| 100 | 200 | 00:00:20 | 100.00% | 0.0101 | 0.0010 |

Training finished: Max epochs completed.

Network training complete.

--> Training bounding box regression models for each object class...100.00%...done.

Detector training complete.

When the model has finished training, we can try with another photo or any Google Maps photo. As we can see, it will get confused very easy.

```
test_image_traffic_lights = imread('imagenes_maps/2.png');
[bboxes, score, label] = detect(rcnn_traffic_lights, test_image_traffic_lights,
'MiniBatchSize', 128);
[score, idx] = max(score);
bbox = bboxes(idx, :);
annotation = sprintf('%s: (Confidence = %f)', 'Traffic light', score);
outputImage = insertObjectAnnotation(test_image_traffic_lights, 'rectangle', bbox,
annotation);
figure;
imshow(outputImage)
```



WRONG WAY SIGN


```
load('./imagenes_jpg/direccion_prohibida/dataset_direccion_prohibida.mat');
dataset_wrong_way = dataset;

I_wrong_way = imread(dataset_wrong_way.paths{1});
I_wrong_way = insertObjectAnnotation(I_wrong_way, 'Rectangle',
dataset_wrong_way.direccion_prohibida{1}, 'Wrong Way sign', 'LineWidth', 8);

figure;
imshow(I_wrong_way);
```



Once we've seen an example image, we can train the net. In this case, we will implement a different algorithm than the propussed one, we will use Adam. Also we will use our NVIDIA GPU to train faster

```
rcnn_wrong_way = trainRCNNObjectDetector(dataset_wrong_way, cifar10Net, options,
'NegativeOverlapRange', [0 0.3], 'PositiveOverlapRange',[0.5 1]);
```

```
*****
Training an R-CNN Object Detector for the following object classes:
```

```
* rois
```


--> Extracting region proposals from 20 training images...done.

--> Training a neural network to classify objects in training data...

Initializing input data normalization.

| Epoch | Iteration | Time Elapsed (hh:mm:ss) | Mini-batch Accuracy | Mini-batch Loss | Base Learning Rate |
|-------|-----------|----------------------------|------------------------|--------------------|-----------------------|
| 1 | 1 | 00:00:00 | 41.41% | 0.7635 | 0.0010 |
| 25 | 50 | 00:00:04 | 100.00% | 2.5319e-05 | 0.0010 |
| 50 | 100 | 00:00:08 | 100.00% | 5.7926e-07 | 0.0010 |
| 75 | 150 | 00:00:11 | 100.00% | 5.9001e-06 | 0.0010 |
| 100 | 200 | 00:00:15 | 100.00% | 0.0003 | 0.0010 |

Training finished: Max epochs completed.

Network training complete.

--> Training bounding box regression models for each object class...100.00%...done.

Detector training complete.

When the model has finished training, we can try with another photo or any Google Maps photo.

```
test_image_wrong_way = imread('./imagenes_maps/4.png');
[bboxes, score, label] = detect(rcnn_wrong_way, test_image_wrong_way,
'MiniBatchSize', 128);
[score, idx] = max(score);
bbox = bboxes(idx, :);
annotation = sprintf('%s: (Confidence = %f)', 'Wrong way sign', score);
outputImage = insertObjectAnnotation(test_image_wrong_way, 'rectangle', bbox,
annotation);
figure;
imshow(outputImage)
```



YIELD RIGHT OF WAY

To prove the power of CNNs, as we detected pedestrian signs, we can detect easily any other 2D patterns in the street. Maybe its a very repetitive topic but it is very useful since autotonomous driving is the future, and this type of recognition could be implemented in this technology. In this case we will detect 'Ceda el paso' signs. In this case a triangular pattern. With some photos taken by us, we create the dataset and then train the net over the Cifar10 one.

```
load('./imagenes_jpg/ceda/dataset_ceda.mat');  
dataset_ceda = dataset;  
  
I_ceda = imread(dataset_ceda.paths{1});
```

```
I_ceda = insertObjectAnnotation(I_ceda, 'Rectangle', dataset_ceda.ceda{1}, 'Yield
sign', 'LineWidth', 8);

figure;
imshow(I_ceda);
```



Then, the training begins.

```
rcnn_ceda = trainRCNNObjectDetector(dataset_ceda, cifar10Net, options,
'NegativeOverlapRange', [0 0.3], 'PositiveOverlapRange',[0.5 1]);
```

```
*****
Training an R-CNN Object Detector for the following object classes:
```

```
* ceda
```

```
--> Extracting region proposals from 8 training images...done.
```

```
--> Training a neural network to classify objects in training data...
```

```
Initializing input data normalization.
```

```
|=====|
| Epoch | Iteration | Time Elapsed | Mini-batch | Mini-batch | Base Learning |
|        |           | (hh:mm:ss)   | Accuracy   | Loss        | Rate          |
```

| | | | | | |
|-----|-----|----------|---------|------------|--------|
| 1 | 1 | 00:00:04 | 64.06% | 0.6541 | 0.0010 |
| 50 | 50 | 00:00:13 | 100.00% | 7.7516e-06 | 0.0010 |
| 100 | 100 | 00:00:20 | 100.00% | 0.0002 | 0.0010 |

Training finished: Max epochs completed.

Network training complete.

--> Training bounding box regression models for each object class...100.00%...done.

Detector training complete.

When the model has finished training, we can try with another photo or any Google Maps photo.

```
test_image_ceda = imread('./imagenes_maps/10.png');
[bboxes, score, label] = detect(rcnn_ceda, test_image_ceda, 'MiniBatchSize', 128);
[score, idx] = max(score);
bbox = bboxes(idx, :);
annotation = sprintf('%s: (Confidence = %f)', 'Yield sign', score);
outputImage = insertObjectAnnotation(test_image_ceda, 'rectangle', bbox,
annotation);
figure;
imshow(outputImage)
```




Detecting more than one type of object:

To detect multiple classes, the usage of more than one network is required. First we set the networks and labels that we want to use.

```
test_image = imread('./imagenes_jpg/direccion_prohibida/3.jpg');  
  
outputImage = test_image;  
  
% Set the minimum score threshold  
min_score_threshold = 0.2;  
  
networks = {rcnn_pedestrian;rcnn_wrong_way};  
labels = ["Pedestrian sign", "Wrong way sign"]  
  
labels = 1x2 string
```

"Pedestrian ... "Wrong way ...

Then, we iterate over each of the networks previously set.

```
for i= 1:numel(networks)
    network = networks{i};
    % Filter out detections below the threshold
    [bboxes, score, label] = detect(network, test_image, 'MiniBatchSize', 128);
    valid_indices = find(score >= min_score_threshold);
    bboxes = bboxes(valid_indices, :);
    score = score(valid_indices);
    label = label(valid_indices);

    objects = numel(valid_indices);

    % Sort the scores
    [scores_sorted, idx_sorted] = sort(score, 'descend');

    for j = 1:objects
        labelidx = label(idx);
        bbox = bboxes(idx_sorted(j), :);
        annotation = sprintf('%s: (Confidence = %f)', labels(i), scores_sorted(j));
        outputImage = insertObjectAnnotation(outputImage, 'rectangle', bbox,
annotation);
    end
end

figure;
imshow(outputImage);
```

